

PARALLEL PROGRAMMING

TOPICS:

- 01. INTRO
- 02. PARALLEL PATTERN 1
- 03. PARALLEL PATTERN 2
- 04. PARALLEL PATTERN 3
- 05. PTHREAD, OPENMP 1
- 06. OPENMP 2
- 07. OPENMP 3
- 08. MPI 1
- 09. MPI 2
- 10. HALIDE
- 11. PARALLEL PATTERNS APPLIED

01. INTRO

MOTIVATION TIME, MONEY, COMPLEX PROBLEMS

- MOOR'S LAW LIMITS POWER CONSUMPTION

PARALLELIZATION:

- AUTOMATIC - NOT FEASIBLE TOOLS CAN'T EXTRACT ALL PARALLELISM
BY HAND TOOLS ONLY COMPILE

TYPES OF PARALLELISM

FLYNN'S TAXONOMY: INSTRUCTION, DATA SISD, MISD, SIMD, MIMD

- LEVEL
 - BIT
 - INSTRUCTION
 - TASK PARALLEL TASK GRAPH, PIPELINE, COMMUNICATION

PARALLEL PROGRAMMING

- NEW PROGRAMMING LANGUAGES 😞 COMPILERS, NEW LANGUAGE
- EXTENSIONS 😊 EASY DESIGN, COMPILER 😞 ONLY SOME TYPE OF PARALLELISM

02. PARALLEL PATTERNS 1

DEPENDENCIES EXECUTION ASSUMPTION, SEQUENTIAL CONSISTENCY

- TRUE (FLOW) DEPENDENCE RAW $out(S_1) \subset in(S_2) \wedge S_1 \text{ d } S_2$
- ANTI-DEPENDENCE WAR $in(S_1) \subset out(S_2) \wedge S_1 \text{ d}^{-1} S_2$
- OUTPUT-DEPENDENCE WAW $out(S_1) \subset out(S_2) \wedge S_1 \text{ d}^0 S_2$
- LOOP-CARRIED DEPENDENCE
 - CAN PREVENT LOOP ITERATION PARALLELIZATION
 - LEXICALLY FORWARD / BACKWARD

PARALLEL PATTERN = RECURRING COMBINATION OF TASK THAT SOLVES A SPECIFIC PROBLEM IN PARALLEL ALGORITHM DESIGN

NESTING PATTERN → ALLOWS OTHER PATTERNS TO BE COMPOSED IN A HIERARCHY

- PATTERN DIAGRAM VISUALIZES PATTERN IDEA

↳ TASK BLOCK = LOCATION OF GENERAL CODE IN AN ALGORITHM CAN BE ANOTHER PATTERN

SERIAL CONTROL PATTERNS: SEQUENCE, SELECTION, ITERATION, RECURSION (+ NESTING FOR COMBINATION)

PARALLEL CONTROL PATTERN EXTEND SERIAL, RELAXED ASSUMPTIONS

- FORK-JOIN FLOW FORKS INTO MULTIPLE PARALLEL FLOWS, THAT REJOIN LATER (SYNC ≠ BARRIER)
- MAP ELEMENTAL FUNCTION PERFORMED OVER EVERY ELEMENT OF A COLLECTION
- STENCIL ELEMENTAL FUNCTION ACCESSES SET OF NEIGHBORS – GENERALIZATION OF MAP
- REDUCTION COMBINER FUNCTION (ASSOCIATIVE) COMBINES EVERY ELEMENT IN A COLLECTION
- SCAN COMPUTES ALL PARTIAL REDUCTIONS OF A COLLECTION
- RECURRENCE COMPLEX VERSION OF MAP (LOOP ITERATION CAN HAVE DEPENDENCIES)

SERIAL DATA MANAGEMENT PATTERNS

- **RANDOM READ/WRITE** POINTERS REFER TO MEMORY ADDRESSES. ALIASING CAN CAUSE PROBLEMS
- **STACK ALLOCATION** FOR DYNAMIC ALLOCATION, ARBITRARY AMOUNT IN CONSTANT TIME. PRESERVES LOCALITY 😊
- **HEAP ALLOCATION** WHEN DATA CAN'T BE ALLOCATED IN LIFO. SLOWER, MORE COMPLEX 😞
- **OBJECTS** LANGUAGE CONSTRUCTS: DATA + CODE TO MANAGE THEM. PART OF CLASS OBJECTS

PARALLEL DATA MANAGEMENT PATTERNS

- **PACK** ELIMINATE UNUSED SPACE IN A COLLECTION + **UNPACK**
- **PIPELINE** PRODUCER-CONSUMER TASK CONNECTION
- **GEOMETRIC DECOMPOSITION** ARRANGE DATA INTO SUBCOLLECTIONS (DON'T MOVE DATA)
- **GATHER** READ COLLECTION OF DATA, GIVEN COLLECTION OF INDICES
- **SCATTER** WRITE OUTPUT AT GIVEN INDEX. POSSIBLE RACE CONDITIONS!
- OTHERS: **SUPERSCALAR SEQUENCES, FUTURES, SPECULATIVE SELECTION, WORKPILE**
SEARCH, SEGMENTATION, EXPAND, CATEGORY REDUCTION

MAP "FOR EACH" LOOP, EACH ITERATION INDEPENDENT \rightsquigarrow CAN BE APPLIED WITHOUT KNOWLEDGE OF NEIGHBORS

- SERIAL vs PARALLEL
- INDEPENDENCE + NO SHARED STATE OTHERWISE · RACES, UNDEFINED BEHAVIOR 😞
- OPTIMIZATIONS
 - SEQUENCES OF MAPS NO WRITE ALL INTERMEDIATE STAGES TO MEM¹
 - CODE FUSION 😞 ARITHMETIC INTENSITY 😊 MEMORY USAGE
 - CACHE FUSION BREAK WORK INTO BLOCKS
- **STENCIL** MAP + NEIGHBORS
- **WORKPILE** WORK ITEMS ADDED TO MAP WHILE IN PROGRESS
- **DIVIDE AND CONQUER** BASE CASE SOLVED SERIALY
- EXAMPLE: SCALED VECTOR ADDITION (SAXPY) $y = \alpha x + y$

03. PARALLEL PATTERNS 2

COLLECTIVES DEAL WITH A COLLECTION AS A WHOLE (REDUCE, SCAN, PARTITION, SCATTER, GATHER)

REDUCE COMBINE COLLECTION OF EL INTO ONE SUMMARY VALUE

- COMBINER FUNCTION ASSOCIATIVE, ELEMENTS PAIRWISE
- VECTORIZATION, TILING
- PRECISION WITH FLOAT 😞
- EXAMPLE: DOT PRODUCT $\text{MAP}(\cdot) + \text{REDUCE}(+)$
- EXAMPLE: MERGE SORT MAP-REDUCE
 - RIGHT BIASED 😞 $O(m^2)$, INSERTION SORT
 - TREE-SHAPED SORT $O(m \log m)$ 😞 NOT A LOT TO PARALLELIZE

SCAN GENERATES NEW SEQUENCE WITH PARTIAL REDUCTIONS

- INCLUSIVE / EXCLUSIVE
- EXAMPLE: RADIX SORT, QUICKSORT
- PARALLELIZATION: UP SWEEP + DOWN SWEEP, 3 PHASE SCAN WITH TILING

04. PARALLEL PATTERNS 3

DATA MOVEMENT LIMITS PERFORMANCE MORE THAN COMPUTATION

- ACROSS MEMORY LOCALITY, DATA ORGANIZATION
- ACROSS NETWORKS NUMBER OF MESSAGES

GATHER READ INPUT COLLECTION AT INPUT INDICES → WRITE OUTPUT COLLECTION

- "MAP + RANDOM READ"
- OUTPUT EL. TYPE = SOURCE EL. TYPE
- OUTPUT SHAPE = INDEX COLLECTION SHAPE

- SPECIAL CASES
 - **SHIFT** + DEFAULT VALUE, DUPLICATE, ROTATE
 - **ZIP** example: REAL+IMAGINARY PARTS → COMPLEX NUMBERS ~> COMBINES IN PAIRS
 - **UNZIP**

SCATTER

- "MAP + RANDOM WRITES"
- COLLISIONS PARALLEL WRITES TO THE SAME LOCATION
L> UNDEFINED RESULT => NEED RULES
- COLLISION RESOLUTION
 - ATOMIC SCATTER NON-DETERMINISTIC
 - PERMUTATION SCATTER ILLEGAL COLLISIONS → OUTPUT PERMUTATION OF INPUT (=> GATHER)
 - MERGE SCATTER ASSOCIATIVE AND COMMUTATIVE OPERATIONS TO MERGE
 - PRIORITY SCATTER PRIORITY BASED ON ITS POSITION

PACK TO ELIMINATE UNUSED ELEMENTS FROM MEM + MOVED TO CONTIGUOUS MEM

- BOOL ARRAY, ~CUMSUM, WRITE OUTPUT WITH OFFSET!

UNPACK

GENERALIZATION OF PACK

- **SPLIT** ELEMENTS MOVED LEFT/RIGHT, NO INFO LOST

UNSPILT

- **BIN** SPLIT WITH MORE CATEGORIES

MAP + PACK

- **EXPAND** EACH EL CAN PRODUCE ANY NUMBER OF EL + RES FUSED TOGETHER

PARALLELIZING ALGORITHMS DIVIDING DATA INTO SECTION

- DIVIDE AND CONQUER
- FORK - JOIN

- GEOMETRIC DECOMPOSITION
- PARTITIONING NON OVERLAPPING, EQUAL-SIZED REGIONS
- SEGMENTATION NON OVERLAPPING, NON-UNIFORM REGIONS

AoS ARRAY OF STRUCTURES

SoA STRUCTURE OF ARRAYS BETTER FOR VECTORIZATION

05. PTHREADS, OPENMP 1

1. DESIGN PARALLEL ALGORITHM UNDERSTAND DATA DEPENDENCIES

2. DESIGN PARALLEL PROGRAM ANALYZE TARGET ARCHITECTURE, CHOOSE LANGUAGE, ANALYZE COMMUNICATIONS

EVALUATION — SEQUENTIAL ALGORITHMS → **TIME, MEMORY COMPLEXITY**
 — PARALLEL ALGORITHMS → **TIME, RESOURCE COMPLEXITY**

QUALITATIVE EVALUATION: DAG, WORK, SPAN, PARALLELISM

PP MODELS SHARED MEMORY, THREADS, MESSAGE PASSING/DISTRIBUTED MEM, DATA PARALLEL

- VERILOG/VHDL 😊 COMMUNICATION+MEM CONTROL, NO OVERHEAD ❌ SPECIFIC HW, DIFFICULT TO LEARN, NO DIFF ARCH
- MPI 😊 DIFFERENT ARCH, SYNCH+DATA COMM EXPLICITLY MANAGED ❌ COMM OVERHEAD, MORE DIFFICULT
- PTHREAD 😊 DIFFERENT ARCH, EXPLICIT PARALLELISM, APP CONTROL ❌ TASK MANAGEMENT CH, NOT SCALABLE, LOW-LEVEL API
- OPENMP 😊 EASY LEARN, SCALABLE ❌ ONLY SH-MEM HOMOGENEOUS SYS, SMALL TASK INTERACTION
- CUDA 😊 GPU, EASY KERNEL WRITING, OPTIMIZED LIBRARIES ❌ ONLY NVIDIA GPU, DIFFICULT MASSIVE PAR.+ KERNEL OPT
- OPENCL 😊 TARGET INDEPENDENT, HIDES ARCH DETAILS ❌ DIFFICULT TO PROGRAM+OBTAIN BEST PERFORMANCE
- APACHE SPARK 😊 API FOR DIFFERENT LANG, NO EXPLICIT PARALL/COMM ❌ ONLY FOR BIG DATA APP, NO GPU FULL SUPPORT
- MIX OPENMP+CUDA, MPI+OPENMP, OPENCL+VERILOG/VHDL

PROCESS INFO ABOUT RESOURCES, EXEC STATE

THREAD INDEPENDENT STREAM OF INSTR WITHIN A PROCESS

- CAN EXECUTE AT THE SAME TIME

- HAS LOCAL RESOURCES, CAN ACCESS SHARED PROCESS RESOURCES
- IMPLICIT COMMUNICATION (SHARED VAR) → EXPLICIT SYNCHRONIZATION TO AVOID DATA RACE

THREADED PROGRAM

- IF ORGANIZABLE IN DISCRETE CONCURRENT TASKS, THREAD SAFE
- MODELS: MANAGER/WORKERS, PIPELINE
- STANDARD IMPLEMENTATIONS: PTHREAD, OPENMP

PTHREADS

- THREAD MANAGEMENT CREATE, JOIN / DETACH, BARRIER INIT / WAIT, CANCEL / EXIT
- SYNCHRONIZATION MUTEX, CONDITION VARIABLES
- examples CALCULATION OF π (MONTECARLO), MULTIRATE BAND-PASS FILTER, PRODUCER-CONSUMER

OPENMP API FOR MULTI-THREADED SHARED-MEM PROGRAMMING

- BASED ON FORK-JOIN MASTER + CONCURRENT SLAVES
- LANG EXTENSIONS: PRAGMA PREPROCESSOR DIRECTIVES, IDENTIFY TASKS. IMPLIED BARRIER AT THE END
 - PARALLEL [NUM THREADS, IF]
 - WORK SHARING: FOR, SECTIONS, SINGLE / MASTER, TASK

06. OPENMP 2

OPENMP LANGUAGE EXTENSIONS

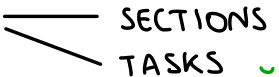
- SYNCHRONIZATION: CRITICAL [NAME], BARRIER, ATOMIC
- DATA ENVIRONMENT: DATA SCOPE ATTRIBUTE CLAUSES (PRIVATE, SHARED, DEFAULT, REDUCTION)
 - SEQUENTIAL MEM CONSISTENCY LOAD/STORE IN ORIGINAL SEQUENTIAL ORDER
 - RELAXED MEM CONSISTENCY EACH THREAD TMP VIEW. FLUSH → TO HAVE SAME GLOBAL VIEW
- RUNTIME FUNCTIONS GET NUM THREADS, GET THREAD NUM, SET NUM THREADS, GET WTIME, GET WTICK
- examples: CALCULATION OF π (INTEGRAL)

NESTED PARALLELISM PARALLEL, WORKSHARING \rightsquigarrow BUT ID RESTARTS FROM 0!

THREAD CANCELLATION CANCEL + CANCELLATION POINT

07. OPENMP 3

OPENMP TASK = BLOCK OF CODE IN A PARALLEL REGION, THAT CAN BE EXECUTED SIMULTANEOUSLY WITH OTHER TASKS IN THE SAME REGION

- PIPELINE  \rightsquigarrow COMMUNICATION AND LOAD IMBALANCE PROBLEM
TASKS \checkmark CAN EXECUTE WHENEVER INPUT IS READY
- TASKWAIT / TASKGROUP
- DEPEND IN/OUT
PRIORITY IMPORTANT TASKS TO BE EXECUTED MORE FREQUENTLY
- TASKLOOP

SIMD VECTORIZATION INSTRUCTIONS ON BLOCK OF DATA (VECTORS)

- COMPILER ISSUES LOOP DEPENDENCIES, POINTER ALIASES, DATA ALIGNMENT, LOOP BOUNDS
- VECTOR REGISTERS LOOP CHUNK SHOULD FIT IT
- simdlen PREFERRED SIZE (PERFORMANCE)

HETEROGENEOUS ARCHITECTURES

- #pragma omp target [map(...)]
 - CODE TO ACCELERATOR, IF PRESENT
 - HOST THREAD BLOCKED
 - map COPIES VARIABLES IN TARGET MEMORY (IT'S SHARED) + to, from, tofrom
- BATCH PROCESSING / COOPERATIVE PATTERNS

08. MPI 1

SHARED MEM ARCHITECTURES 😞 HARD TO BUILD 😊 EASY TO PROGRAM (OPENMP, C++ LIBRARY)

DISTRIBUTED MEM ARCHITECTURE REQUIRES EXPLICIT COMMUNICATION 😞 SYNCHRONIZATION

MPI (MESSAGE PASSING INTERFACE) STANDARD THAT DEFINES INTERFACES FOR DISTRIBUTED MEM COMMUNICATIONS

- OUTCOME: PDF DOCUMENT
- IMPLEMENTATIONS OPENMPI, MPICH
 - libmpi.so IMPLEMENTATION LIBRARY
 - mpicc COMPILER
 - mpiexec LAUNCHER
 - mpirun FORK+EXEC CODE
- INIT + THREAD USAGE
- FINALIZE
- COMMUNICATOR CONTEXT

POINT TO POINT COMMUNICATION

- SEND, RECV(STATUS), PROBE
- WAIT, TEST, CANCEL

09. MPI 2

COLLECTIVE COMMUNICATIONS

- COMMUNICATOR HANDLING DUP, SPLIT, FREE
- SYNCHRONIZATION BARRIER
- DATA TRANSFER BROADCAST, GATHER, SCATTER
- REDUCTIONS REDUCE (MAX, SUM, LAND, MAXLOC)

10 HALIDE

POWER = OP / SECONDS · JOULES / OP

GENERAL-PURPOSE PROCESSORS NOT ENERGY EFFICIENT

HETEROGENEOUS PROCESSING

- GPU
- DSP DIGITAL SIGNAL PROCESSORS
- ASIC DOMAIN-SPECIFIC HW EXTREME SPECIALIZATION
- FPGA

CHALLENGES — SW DESIGN ALGO THAT CAN BE DECOMPOSED AND MAP WELL ON HW
— HW CHOOSE RESOURCES, CHIP AREA FOR EACH FUNCTIONALITY

ENERGY-EFFICIENT COMPUTING REDUCE COMPUTATION, DATA TRANSFER, CHOOSE BEST COMPUTE UNIT

DSL PRODUCTIVITY + PERFORMANCE, NO GENERALITY

IMAGE PROCESSING DEMANDS OPTIMIZATION → FASTER ALGO, HW, MORE EFFICIENT USE OF HW

- C++ WITH MULTITHREADING
- CUDA / OPENCL
- OPTIMIZED LIBRARIES

HALIDE DSL EMBEDDED IN C++ + SPECIALIZED COMPILER

- DECOUPLE ALGO FROM SCHEDULE
- ☹️ ALL COMPUTATIONS OVER REGULAR GRIDS, MANUALLY SPECIFY SCHEDULE (AUTOSCHEDULER CAN SUGGEST)
- CAN OPTIMIZE WITH
 - ORDER OF VALUES COMPUTED SERIAL, VECTORIZE, TILING
 - WHEN COMPUTE RESULTS eg: COMPUTE ALL/SUBSET OF PRODUCER, THEN PASS TO CONSUMER
- SCHEDULING ACROSS STAGES DRIVEN BY CONSUMER
- examples: BLUR, BRIGHTEN

11. PARALLEL PATTERNS APPLIED

PP MODELS OFFER API FOR PATTERNS

- MAP `#pragma omp parallel for`
- REDUCTION `#pragma omp for reduction (+ : var)`
- WORKPILE example: TREE SEARCH
- SCAN + 3 PHASE TILING APPROACH

+ SORTING